

RYO KAMOI

ryokamoi@psu.edu

<https://ryokamoi.github.io/>**Research Interests**

- Self-Evolution of LLMs and LLM Agents [[TACL'24](#), [COLM'24](#)]
- Verification of Intermediate Steps of LLMs and LLM Agents: Verification of LLM Reasoning [[ACL'26 Findings](#)]
- Evaluation of Open-Ended Non-Verifiable Tasks: Evaluation of Simulated Conversations [[EMNLP'26](#)], Faithfulness Evaluation [[EMNLP'23](#), [EACL'23](#)]
- Vision-Language Models [[COLM'25](#)]

Education

Pennsylvania State University - Ph.D. Candidate in Computer Science Advisor: Dr. Rui Zhang	Aug 2023 -
University of Texas at Austin - Master of Science in Computer Science Advisor: Dr. Greg Durrett, Mentor: Dr. Tanya Goyal	Aug 2020 - Dec 2022
Keio University , Japan - Bachelor of Engineering in Statistics Advisor: Dr. Kei Kobayashi	Apr 2016 - Mar 2020

Research Internships

Google , Sunnyvale, CA – Student Researcher at the Cloud AI Research team	Jun - Sep 2026
Microsoft , Redmond, WA – Research Intern at the Office of Applied Research	May - Dec 2025
Amazon , Cambridge, UK – Applied Scientist Intern at the Alexa team	Jul - Dec 2021

Selected Papers<https://scholar.google.com/citations?user=4OWTLKAAAAAJ>

- Ryo Kamoi**, Ameya Godbole, Binglin Zhou, Xiaoxin Lu, Longqi Yang, Rui Zhang, Mengting Wan, Pei Zhou. 2026. [Evaluating LLM-Simulated Conversations in Modeling Inconsistent and Uncollaborative Behaviors in Human Social Interaction](#). *EMNLP 2026*.
- Ryo Kamoi**, Yusen Zhang, Nan Zhang, Sarkar Snigdha Sarathi Das, Ranran Haoran Zhang, Wenpeng Yin, Rui Zhang. 2026. [Efficient PRM Training Data Synthesis via Formal Verification](#). *ACL 2026 Findings*.
- Ryo Kamoi**, Yusen Zhang, Sarkar Snigdha Sarathi Das, Ranran Haoran Zhang, Rui Zhang. 2025. [VisOnlyQA: Large Vision Language Models Still Struggle with Visual Perception of Geometric Information](#). *COLM 2025*.
- Ryo Kamoi**, Yusen Zhang, Nan Zhang, Jiawei Han, Rui Zhang. 2024. [When Can LLMs Actually Correct Their Own Mistakes? A Critical Survey of Self-Correction of LLMs](#). *TACL 2024*.
- Ryo Kamoi**, Sarkar Snigdha Sarathi Das, Renze Lou, Jihyun Janice Ahn, Yilun Zhao, Xiaoxin Lu, Nan Zhang, Yusen Zhang, Ranran Haoran Zhang, Sujeeth Reddy Vummanthala, Salika Dave, Shaobo Qin, Arman Cohan, Wenpeng Yin, Rui Zhang. 2024. [Evaluating LLMs at Detecting Errors in LLM Responses](#). *COLM 2024*.
- Ryo Kamoi**, Tanya Goyal, Juan Diego Rodriguez, and Greg Durrett. 2023. [WiCE: Real-world Entailment for Claims in Wikipedia](#). *EMNLP 2023 (main)*.
- Ryo Kamoi**, Tanya Goyal, Juan Diego Rodriguez, and Greg Durrett. 2023. [Shortcomings of Question Answering Based Factuality Frameworks for Error Localization](#). *EACL 2023 (main)*.

Media Mentions

Self-Correction in Large Language Models. **Communications of the ACM.**

Feb 26, 2025

Honors and Awards

Research Papers

First Place in the Poster Session at industryXchange 2026 2026

Best Paper Award at the 2nd AI4Research workshop @ AAAI 2025 2025

Services

Great Reviewer at ACL 2026 (133/8,594) 2026

Top 200 Reviewer at ICLR 2026 (200/18,054) 2026

Outstanding Reviewer at EMNLP 2024 2024

Scholarships & Academic Honors

2026-27 Herbert J. Bailey Graduate Fellowship, Penn State University 2026

2025-26 Vice Provost and Dean of the Graduate School Student Persistence Scholarship, Penn State University 2026

Keio University Global Fellowship - Scholarships for my graduate study at UT Austin 2020

Keio Engineering Foundation Award - Graduation with highest honors (First place in the Dept. of Mathematics) 2020

Japan Student Services Organization (JASSO) Exchange Student Scholarship 2018

Academic Services

Reviewer (*outstanding reviewer)

ARR (2023-present), COLM 2024-present, **ACL 2026***, **ICLR 2026***, NeurIPS 2025, ICCV 2025, CVPR 2025, **EMNLP 2024***, AISTATS 2021, BMVC 2020

Conferences and Workshops - MASC-SLL 2025 (organizer), EMNLP 2024 (volunteer)

Events - NLP Colloquium JP (organizer, 2024 -)

Educational Experience

Mentorship

Avitej Iyer (M.S. Student at Penn State University) 2025

Teaching Assistant, Penn State University

CMPSC 442: Artificial Intelligence Spring 2024

CMPSC 448: Machine Learning and AI Fall 2023, Fall 2024, Fall 2025

Invited Talks

Towards Self-Evolving LLMs

AiTech Seminar at Okazaki Laboratory, Institute of Science Tokyo 2025

Michinoku Communication Science Seminar (MiCS) at Tohoku University 2025

AI for Math Minisymposium at MSIAM-NNP 2025

Emergence and Limitations in Self-Refinement of LLMs

Matuo-Iwasawa Lab at the University of Tokyo 2024

Other Papers

First-Author Papers

- Ryo Kamoi** and Kaname Tomite. 2021. [Efficient Unknown Object Detection with Discrepancy Networks for Semantic Segmentation](#). *In the NeurIPS 2021 Workshop on Machine Learning for Autonomous Driving*.
- Ryo Kamoi** and Kei Kobayashi. 2020. [Why is the Mahalanobis Distance Effective for Anomaly Detection?](#) *arXiv preprint arXiv:2003.00402*.
- Ryo Kamoi** and Kei Kobayashi. 2020. [Out-of-Distribution Detection with Likelihoods Assigned by Deep Generative Models Using Multimodal Prior Distributions](#). *In The AAAI's Workshop on Artificial Intelligence Safety*.
- Ryo Kamoi**, Kei Kobayashi. 2019. [Likelihood Assignment for Out-of-Distribution Inputs in Deep Generative Models is Sensitive to Prior Distribution Choice](#). *arXiv preprint arXiv:1911.06515*.

Co-Authored Papers

- Binglin Zhou, Peng Shi, **Ryo Kamoi**, Nan Zhang, Rui Zhang. 2026. [ToolSciVer: Multimodal Scientific Claim Verification with Visual Tool Augmented Reinforcement Learning](#). *arXiv preprint arXiv:2607.16131*.
- Jihyun Janice Ahn, **Ryo Kamoi**, Berk Atıl, Renze Lou, Wonwoo Kang, Heehyun Park, Sarkar Snigdha Sarathi Das, Zhuoyang Zou, Xiaoxin Lu, Yusen Zhang, Asfahan Shah, Ridwanul Hasan Tanvir, Lingxiao Zhao, Hongxi Huang, Vignesh Venkatesh, Dianjun Lin, Hamid Shah, Wentao Wang, Zhanpeng Song, Joshua Reed Bassin, Dax Patel, Ishan Appareddy Agrahar, Sahil Pardasani, Xin Dong, Fatemeh Rahbari, Benjamin David Rishel, Soochan Andrew Lee, Yuv Boghani, Ali B. AlNaseeb, Pranav Suby, Seokhyeon Bae, Shreya Buddharaju, Damien Kula, Soumyadeep Das, Hanyang Frank Liu, Faye Fei Mo, Rui Zhang, Wenpeng Yin. 2026. [Bridging the Know-Act Gap via Task-Level Autoregressive Reasoning](#). *arXiv preprint arXiv:2603.22619*.
- Bowen Jiang, Taiwei Shi, **Ryo Kamoi**, Yuan Yuan, Camillo J Taylor, Longqi Yang, Pei Zhou, Sihao Chen. 2026. [One Model, All Roles: Multi-Turn, Multi-Agent Self-Play Reinforcement Learning for Conversational Social Intelligence](#). *arXiv preprint arXiv:2602.03109*.
- Yusen Zhang, Wenliang Zheng, Aashrith Madasu, Peng Shi, **Ryo Kamoi**, Hao Zhou, Zhuoyang Zou, Shu Zhao, Sarkar Snigdha Sarathi Das, Vipul Gupta, Xiaoxin Lu, Nan Zhang, Ranran Haoran Zhang, Avitej Iyer, Renze Lou, Wenpeng Yin, Rui Zhang. 2025. [HRScene: How Far Are VLMs from Effective High-Resolution Image Understanding?](#) *ICCV 2025*.
- Renze Lou, Hanzi Xu, Sijia Wang, Jiangshu Du, **Ryo Kamoi**, Xiaoxin Lu, Jian Xie, Yuxuan Sun, Yusen Zhang, Jihyun Janice Ahn, Hongchao Fang, Zhuoyang Zou, Wenchao Ma, Xi Li, Kai Zhang, Congying Xia, Lifu Huang, Wenpeng Yin. 2025. [AAAR-1.0: Assessing AI's Potential to Assist Research](#). *ICML, Best Paper Award at the 2nd AI4Research workshop @ AAAI 2025*.
- Sarkar Snigdha Sarathi Das, **Ryo Kamoi**, Bo Pang, Yusen Zhang, Caiming Xiong, Rui Zhang. 2025. [GReaTer: Gradients over Reasoning Makes Smaller Language Models Strong Prompt Optimizers](#). *ICLR 2025*.
- Jihyun Janice Ahn, **Ryo Kamoi**, Lu Cheng, Rui Zhang, Wenpeng Yin. 2024. [Direct-Inverse Prompting: Analyzing LLMs' Discriminative Capacity in Self-Improving Generation](#). *arXiv preprint arXiv:2407.11017*.
- Yilun Zhao, Yitao Long, Hongjun Liu, Linyong Nan, Lyuhao Chen, **Ryo Kamoi**, Yixin Liu, Xiangru Tang, Rui Zhang, Arman Cohan. 2023. [DocMath-Eval: Evaluating Numerical Reasoning Capabilities of LLMs in Understanding Long Documents with Tabular Data](#). *ACL 2024*.

Yusen Zhang, Nan Zhang, Yixin Liu, Alexander Fabbri, Junru Liu, **Ryo Kamoi**, Xiaoxin Lu, Caiming Xiong, Jieyu Zhao, Dragomir Radev, Kathleen McKeown, Rui Zhang. 2024. [Fair Abstractive Summarization of Diverse Perspectives](#). *NAACL 2024*.